

Deep reinforcement learning using sparse distributed memory

Michal CHOVANEC

October 2017

Overview

- Reinforcement learning
- Sparse representation
- Advanced sparse distributed memory
- Experimental results

Reinforcement learning

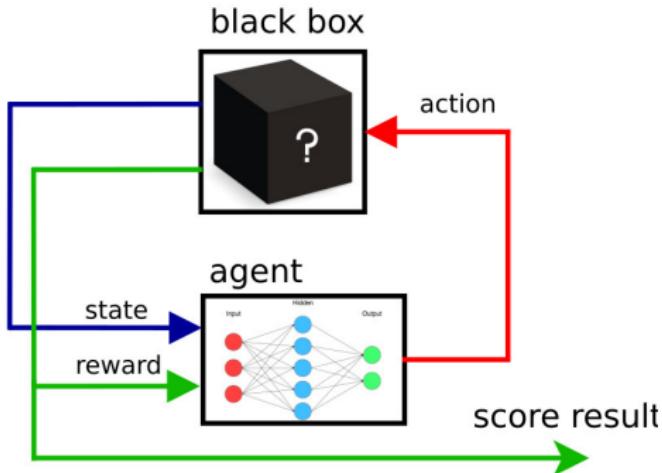
Supervised vs. Reinforcement learning



	input	required output	output
supervised	yes	yes	class
reinforcement	yes	unknown	action

Reinforcement learning

Learning based on rewards and punishments



- input : state, reward
- output: action

SARSA algorithm

State Action Reward State Action ¹

$$Q'(s, a) = (1 - \alpha)Q(s, a) + \alpha(R(s, a) + \gamma Q(s', a'))$$

where

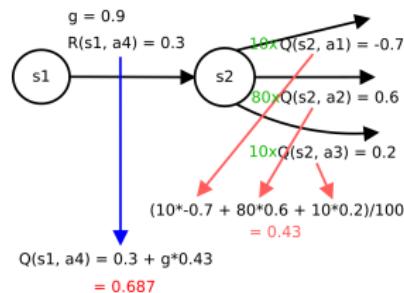
$Q(s, a)$ is previous state

$Q(s', a')$ is actual state

$R(s, a)$ is reward obtained in state s after executing action a

γ is discount factor $\gamma \in \langle 0, 1 \rangle$

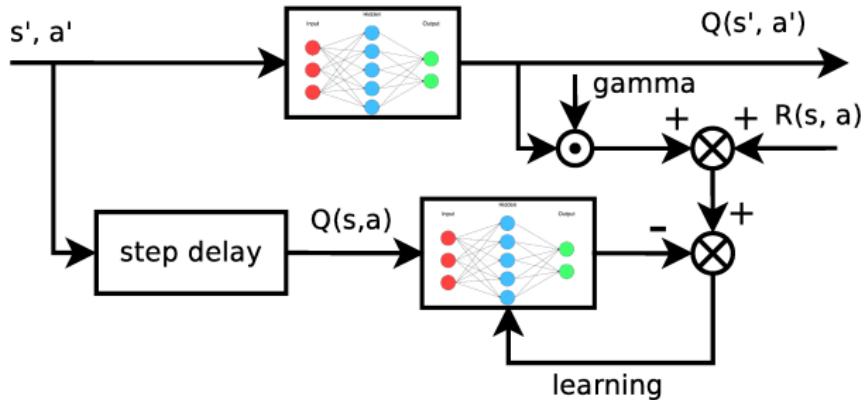
α is learning rate $\alpha \in (0, 1)$



¹ Online Q-Learning using Connectionist Systems, by Rummery & Niranjan (1994)

SARSA algorithm is recurrent

Main learning mechanism :
learn previous $Q(s, a)$ from actual $Q(s', a')$.



SARSA algorithm is recurrent

Main learning mechanism :

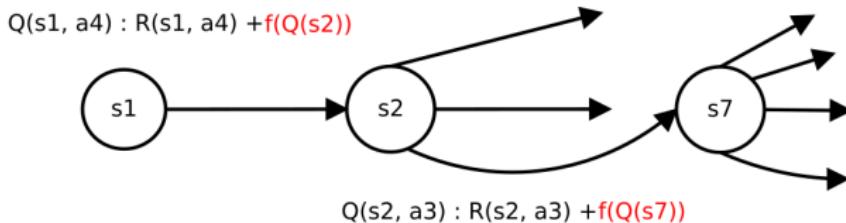
learn previous $Q(s, a)$ from actual $Q(s', a')$.

$$Q(n+1) = R(n+1) + \gamma Q(n+2)$$

$$Q(n) = R(n) + \gamma Q(n+1)$$

$$Q(n-1) = R(n-1) + \gamma Q(n) // \text{ previous slide }^2$$

$$Q(n-2) = R(n-2) + \gamma Q(n-1)$$



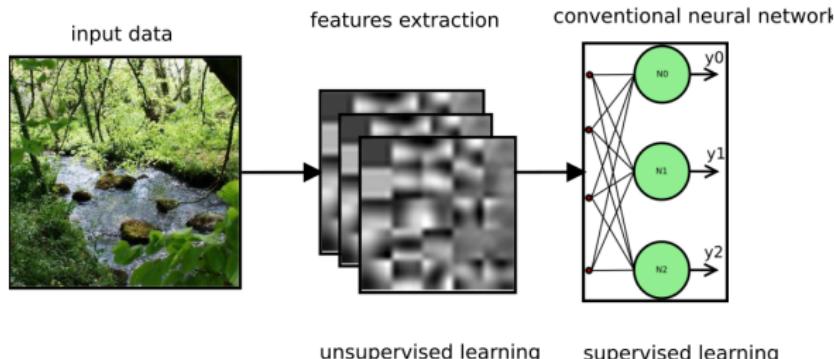
¹ equations are simplified, without s and a

Q value approximation

$S =$



- table
- linear combination of features
- neural network
- sparse representation



Q value approximation

Linear combination of K learned features $f_k(s)$, vector v should be **sparse**, why ?

$$F(s) = \sum_{k=0}^{K-1} v_k f_k(s)$$

The reason for conjecturing that sparseness is an appropriate prior for the a_i is based on the intuition that natural images may generally be described in terms of a small number of structural primitives—for example, edges, lines, or other elementary features (Field, 1994). In addition, one can see evidence for sparse structure in images by filtering them with a set of log-Gabor filters and collecting histograms of the resulting output distributions; these distributions typically show high kurtosis (Field, 1993), which is indicative of sparse structure.

¹ Olshausen and Field (1997): Sparse coding with an overcomplete basis set

² Olshausen BA, Field DJ (2004) : Sparse coding of sensory inputs. Current Opinion in Neurobiology, 14, 481-487

³ Mushroom body, locust (Laurent)

⁴ HVC, zebra finch (Fee)

⁵ Auditory cortex, mouse (DeWeese & Zador)

⁶ Hippocampus, ratprimate(Thompson & Best; Skaggs)

⁷ Motor cortex, rabbit (Swadlow)

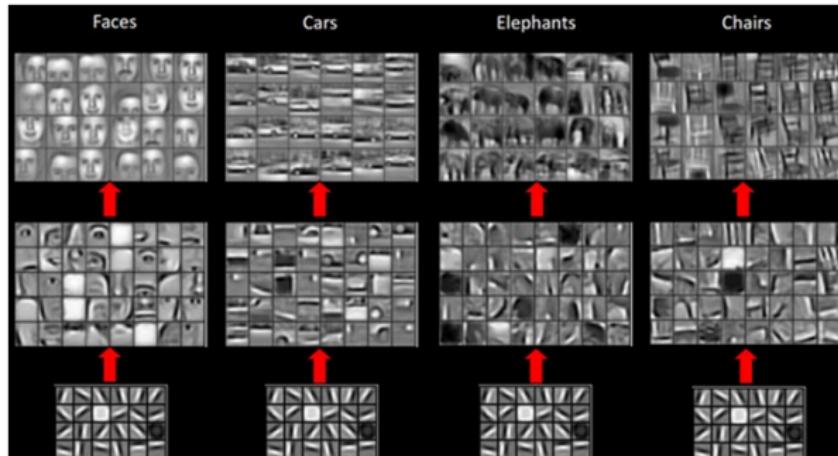
⁸ Visual cortex, monkeycat (Vinje & Gallant)

⁹ Inferotemporal cortex, human (Fried & Koch)

Looking for symmetry

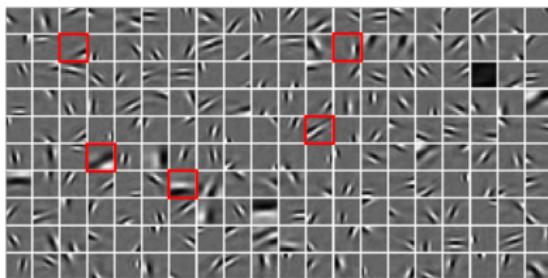
- Euler-Lagrange equation
- Noether's theorem
- Conservation law

$$S(q) = \int_a^b L\left(t, q(t), \frac{dq(t)}{dt}\right) dt$$



Sparse representation

$$F(s) = \sum_{k=0}^{K-1} v_k f_k(s)$$



$$\text{[image]} = 0.4 \text{ [block]} + 0.7 \text{ [block]} + 0.1 \text{ [block]} + 0.2 \text{ [block]} + 0.5 \text{ [block]}$$

Sparse representation

$$E(w, y) = \sum_{k=0}^{K-1} \|x_k - wy_k\|_2^2 + \lambda_0 \|y_k\|_1^2 + \lambda_1 \|w\|_{2,1}^2$$

subject to $y_{ki} > 0$

- reconstruction condition

$$\|x_k - wy_k\|_2^2$$

- y_k sparsity condition

$$\lambda_0 \|y_k\|_1^2$$

- w sparsity condition

$$\lambda_1 \|w\|_{2,1}^2$$

Estimating W and Y - modified K-SVD algorithm

Data: W , depth, X

Result: sparse representation Y , features matrix weights W
initialization

$W \leftarrow \text{random}$, $R \leftarrow X$

η learning rate, ϵ shrink value

for j from 0 to depth do

$$\text{best} = \arg \max_i W[i]R$$

$$d = \text{ReLU}\left(\frac{W[\text{best}]R}{\|W[\text{best}]\|_2^2}\right)$$

$$Y[\text{best}] = d$$

for i from 0 to size(R) do

$$dw = d(R[i] - W[\text{best}][i])$$

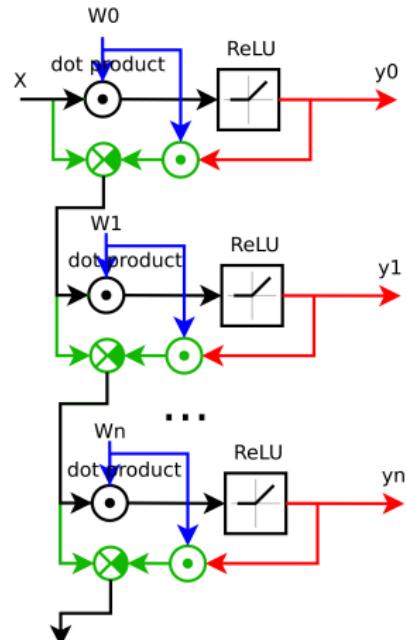
$$W[\text{best}][i] = W[\text{best}][i] + \eta dw$$

$$W[\text{best}][i] = \text{shrink}(W[\text{best}][i], \epsilon)$$

$$R[i] = R[i] - d * W[\text{best}][i]$$

end

end

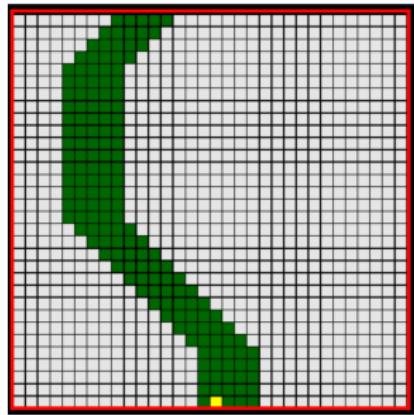


Advanced sparse distributed memory

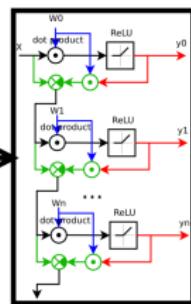
input 1x32x32

sdm 32 32 256 10

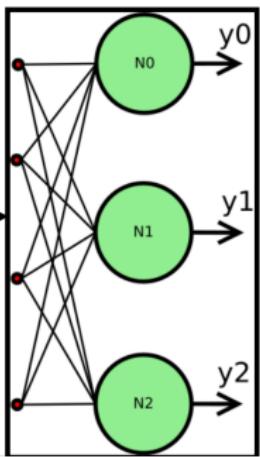
linear layer



1024



256

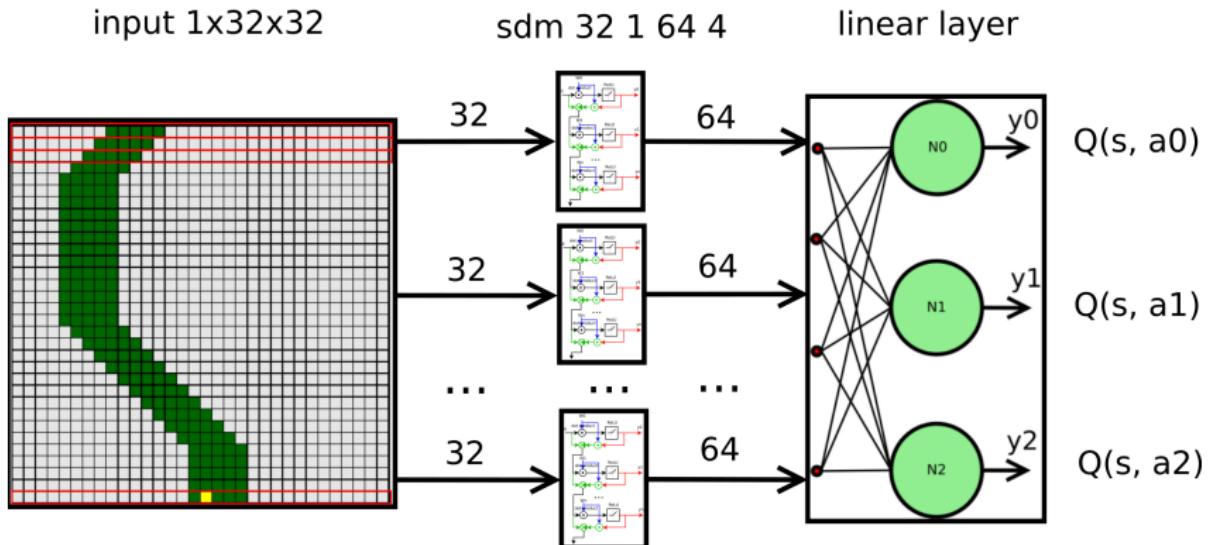


$Q(s, a_0)$

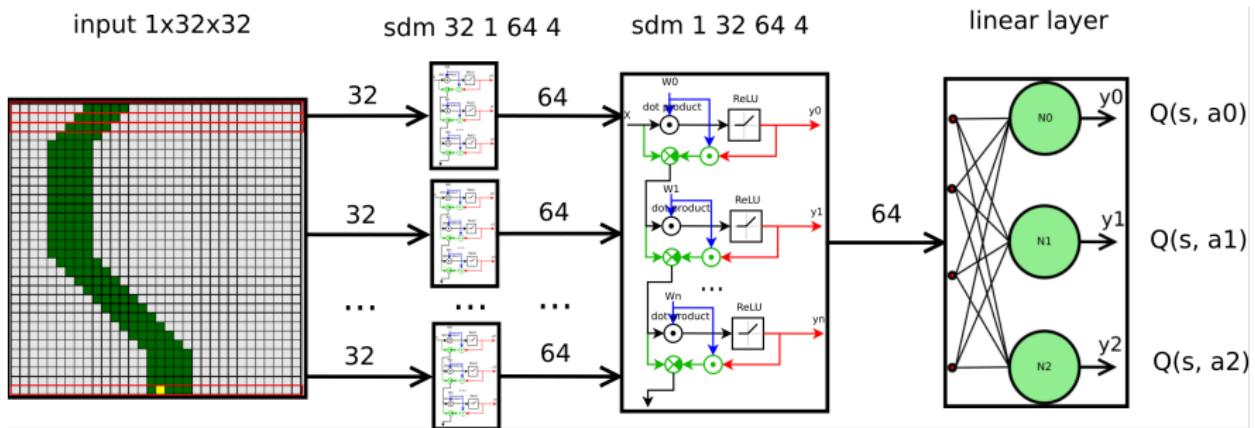
$Q(s, a_1)$

$Q(s, a_2)$

Advanced sparse distributed memory + Convolution



Advanced sparse distributed memory + Convolution + Deep layer

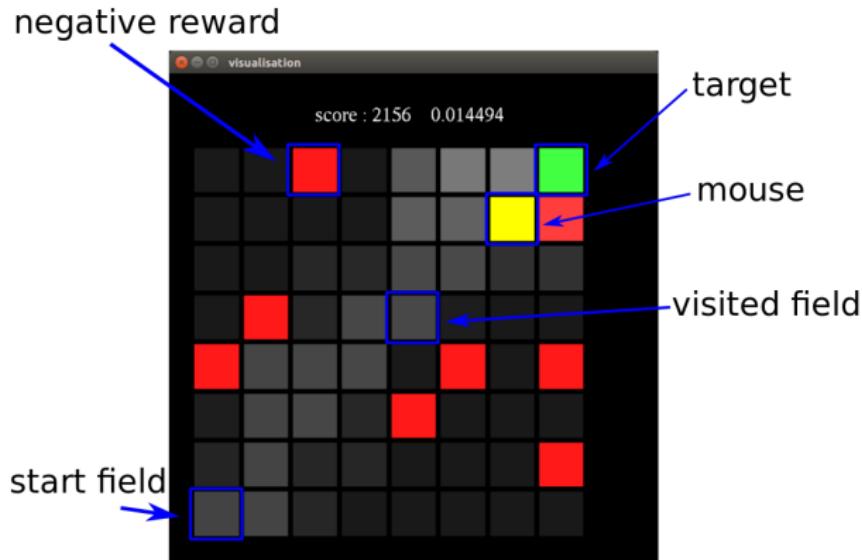


Experimental results

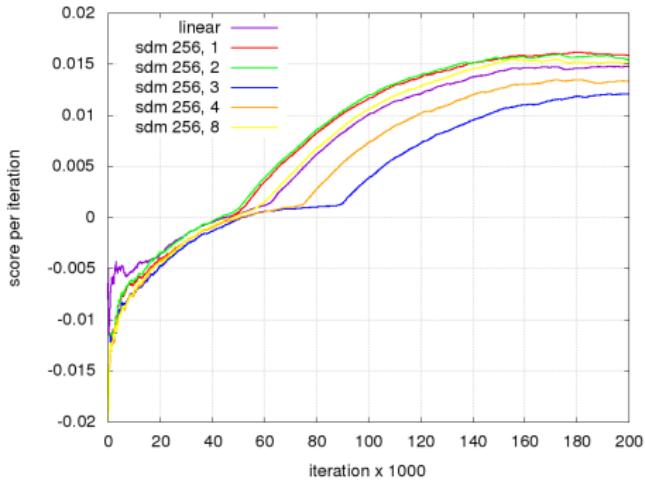
- Mouse and food
- Worms
- Line follower

Mouse and food

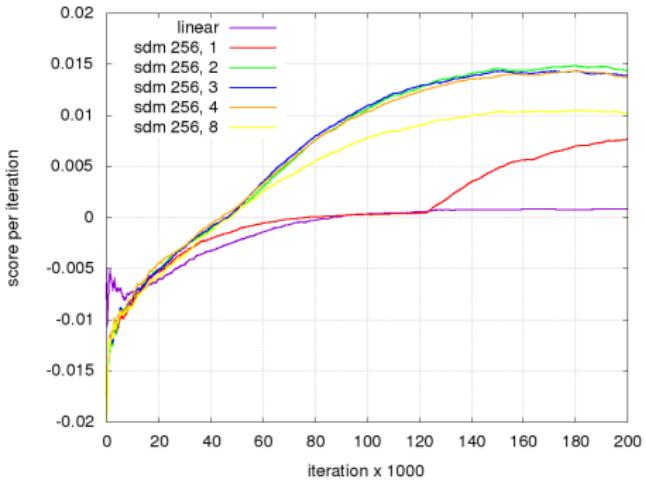
- state : vector $64 \times \{0, 1\}$, 1 if mouse on field
- action : left, right, up, down
- reward : +1 for target hit, -1 for red hit



Mouse and food

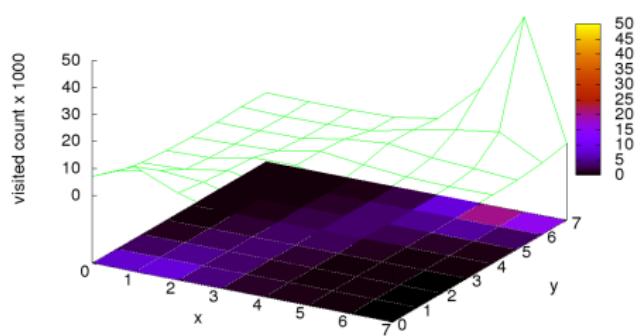


Progress rate, state noise 0.0

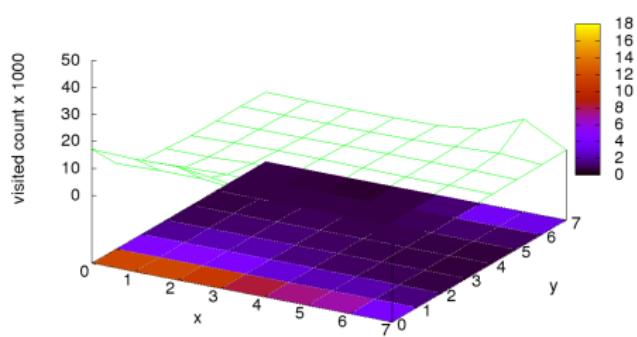


Progress rate, state noise 0.1

Mouse and food - linear approximator

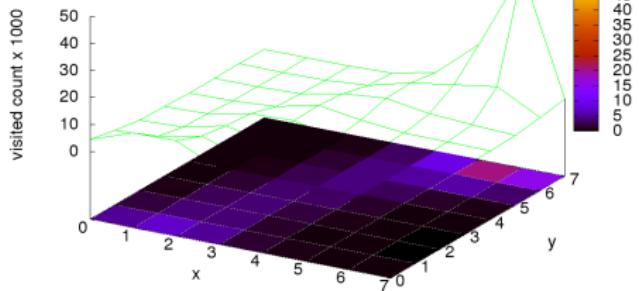


Visited field rate, state noise 0.0

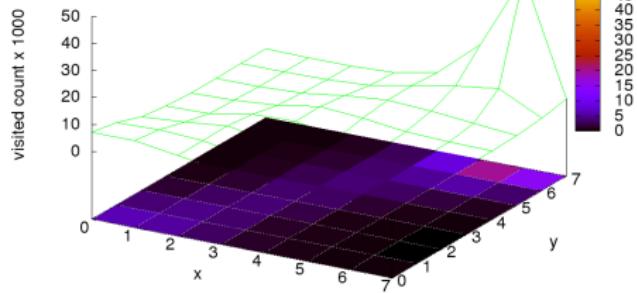


Visited field rate, state noise 0.1

Mouse and food - sdm approximator 256, 2

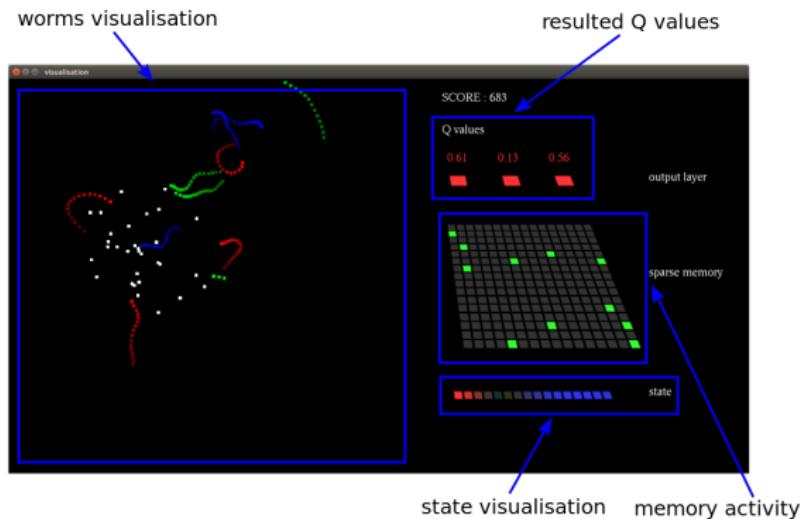


Visited field rate, state noise 0.0

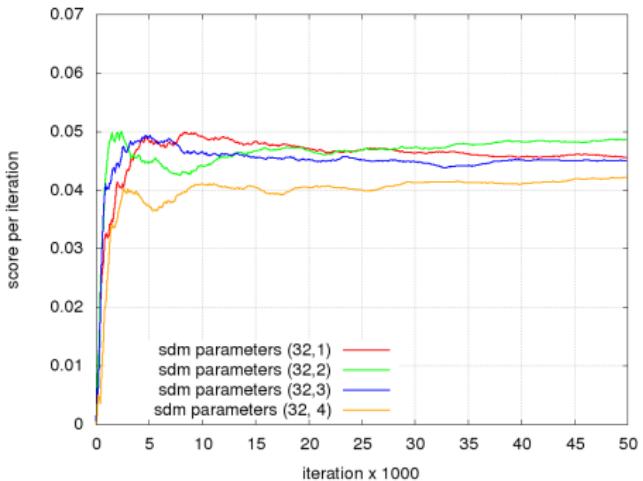
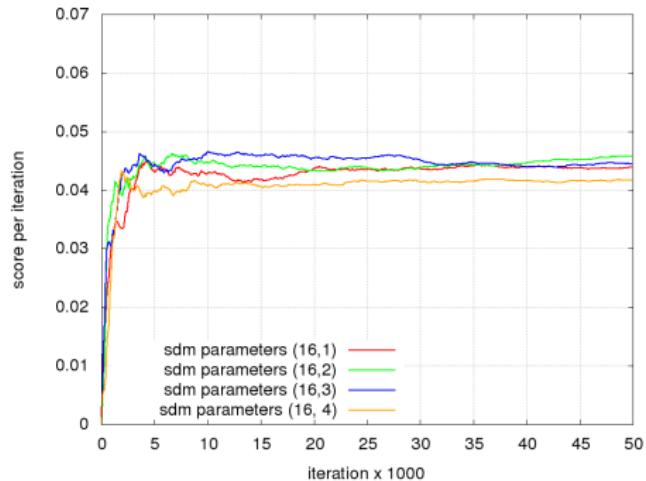


Visited field rate, state noise 0.1

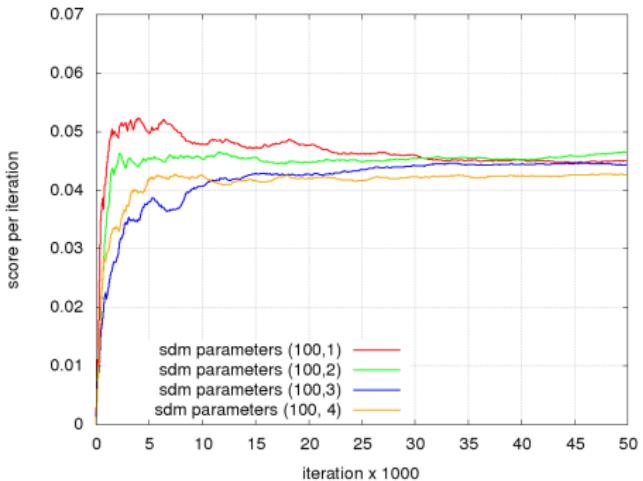
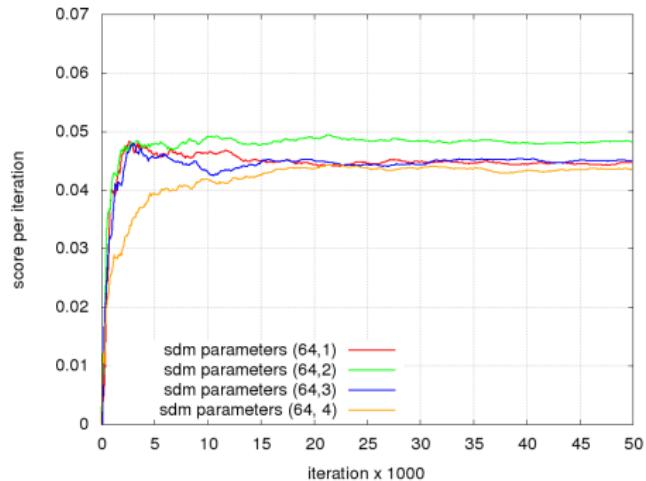
Worms



Worms - memory size and activity results



Worms - memory size and activity results

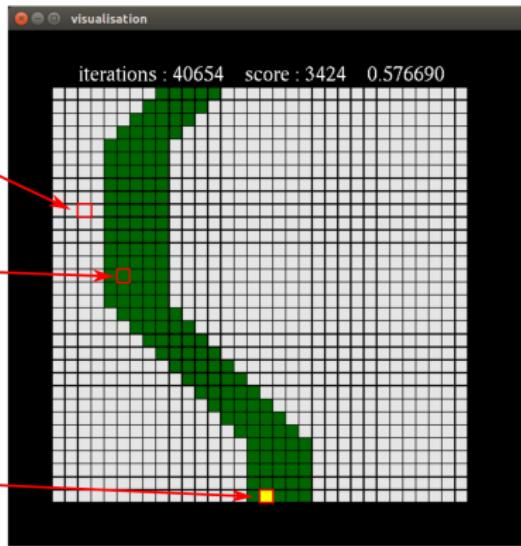


Line follower

negative reward

positive reward

agent



Line follower

- motion equations of robot

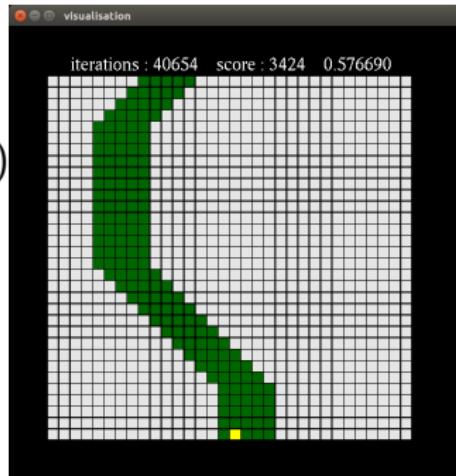
$$a(n) = \{-1, 0, 1\}$$

$$v(n) = \alpha r(n-1) + (1.0 - \alpha)a(n)$$

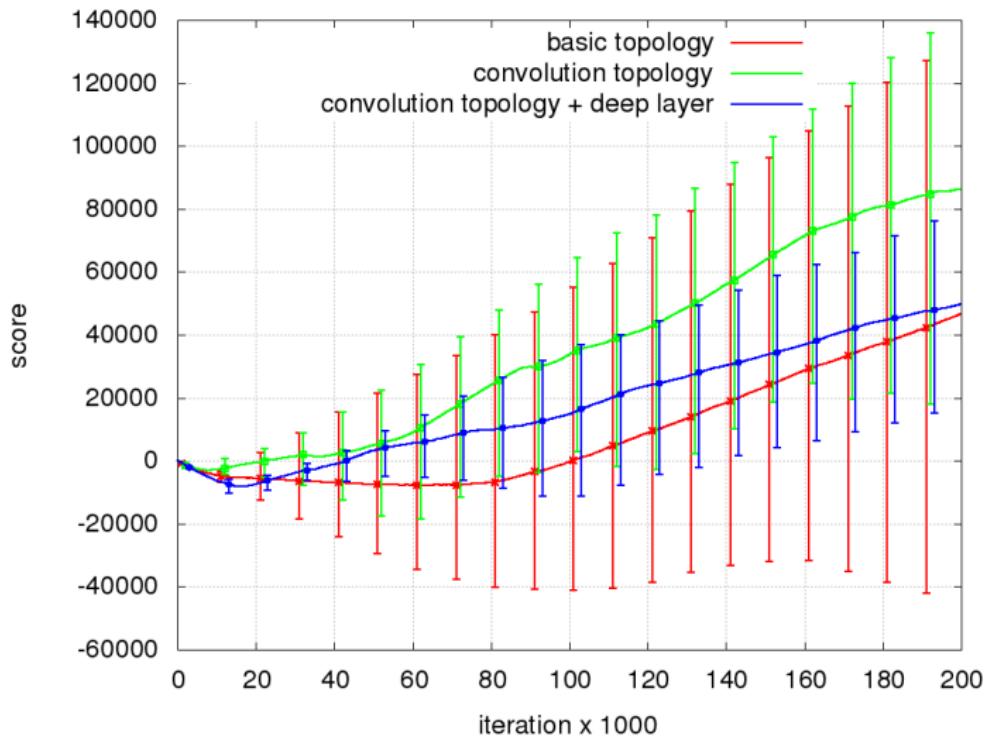
$$p(n) = p(n-1) + v(n)$$

- reward

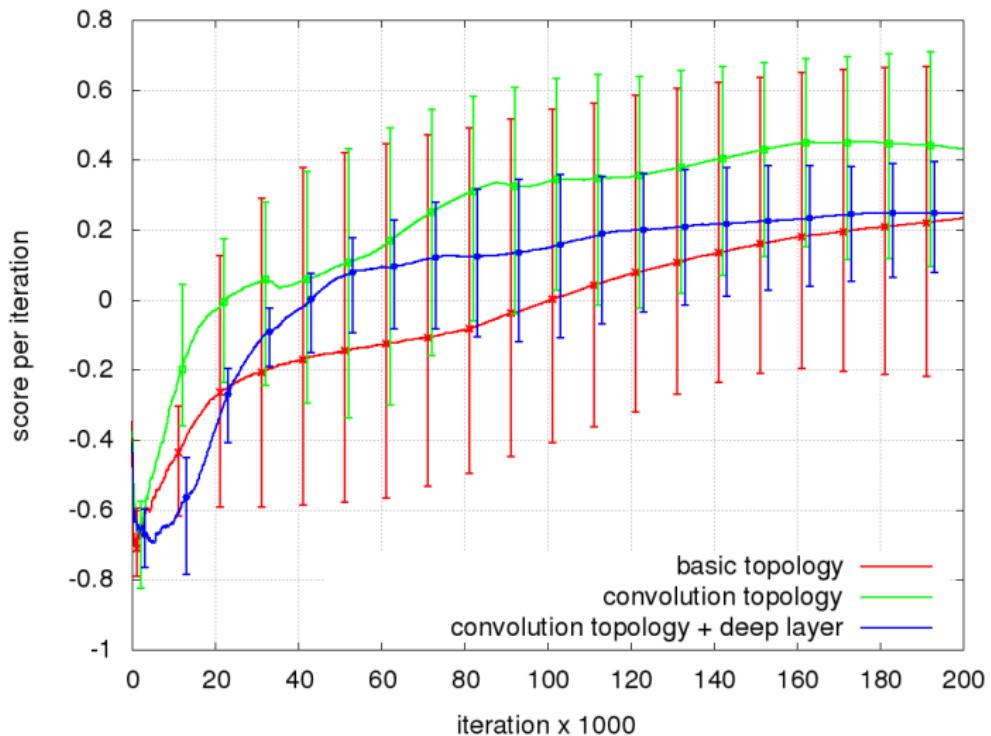
$$r(n) = \begin{cases} +1, & \text{if green field} \\ -1, & \text{otherwise} \end{cases}$$



Line follower - results



Line follower - results



Q&A



<https://github.com/michalnand/robotics>

https://github.com/michalnand/machine_learning_new
michal.nand@gmail.com